

NIE-VCC Tutorial 5.

1. Build an image with a web application using the Dockerfile method.

I created an image which for a web application which shows “Hola {name}!” depending on the given URL (This work is taken from the course NIE-AM2, most specifically the homework 3) the needed files are on the folder “application from another subject”.

Let’s upload the folder to the virtual machine by sftp.

```
sftp> put -r app
Uploading app/ to /home/scast/app/
Entering app/
app/.keep                                100%   0    0.0KB/s   00:00
app/app.js                              100% 283   18.1KB/s   00:00
app/Dockerfile                          100%  68    4.1KB/s   00:00
```

Now let’s install the node:alpine3.15 image because it is a simple image which consumes not so many resources and it is the image in which our server is created.

```
root@tutorial-5:~# docker pull node:alpine3.15
alpine3.15: Pulling from library/node
df9b9388f04a: Pull complete
b81098a29065: Pull complete
d18e917c779e: Pull complete
8c3de7fd67d7: Pull complete
Digest: sha256:0677e437543d10f6cb050d92c792a14e5eb84340e3d5b4c25a88baa723d8a4ae
Status: Downloaded newer image for node:alpine3.15
docker.io/library/node:alpine3.15
root@tutorial-5:~# docker images
REPOSITORY    TAG       IMAGE ID       CREATED        SIZE
node          alpine3.15  9f58095cfeb6   4 days ago    172MB
root@tutorial-5:~#
```

Now inside our app folder let’s execute the command “docker build -t server .” to create our new image which will have the implementation of the server.



```

root@tutorial-5:~# ls
app snap
root@tutorial-5:~# cd app/
root@tutorial-5:~/app# ls
app.js Dockerfile node_modules package.json package-lock.json
root@tutorial-5:~/app# docker build -t server .
Sending build context to Docker daemon 2.043MB
Step 1/4 : FROM node:lts-alpine3.15
node:lts-alpine3.15: Pulling from library/node
df9b9388f04a: Already exists
70c90f7de7cb: Pull complete
f83937c3ce37: Pull complete
98b78bba1d70: Pull complete
Digest: sha256:1a9a71ea86aad332aa7740316d4111ee1bd4e890df47d3b5eff3e5bde3b3d10
Status: Downloaded newer image for node:lts-alpine3.15
--> e5065cc78074
Step 2/4 : COPY . /app
--> b3b925064337
Step 3/4 : WORKDIR /app
--> Running in 3cf8f027c3ac
Removing intermediate container 3cf8f027c3ac
--> 0d59debffc86
Step 4/4 : CMD node app.js
--> Running in 7611ece01481
Removing intermediate container 7611ece01481
--> d6be64a63ab1
Successfully built d6be64a63ab1
Successfully tagged server:latest
root@tutorial-5:~/app#

```

Find our new docker image.

```

root@tutorial-5:~/app# docker images

```

REPOSITORY	TAG	IMAGE ID	CREATED	SIZE
server	latest	d6be64a63ab1	47 seconds ago	113MB
node	alpine3.15	9f58095cfeb6	4 days ago	172MB
node	lts-alpine3.15	e5065cc78074	3 weeks ago	112MB

Let's build our server and create a container which will run it.

The screenshot shows a web browser window on the left with the address bar displaying '10.38.6.236:8080/john' and the page content 'Hello John'. On the right, a terminal window shows the following commands and output:

```

root@tutorial-5:~/app# docker run --name hi_server -p 8080:8080 -d server
root@tutorial-5:~/app# docker ps
CONTAINER ID   IMAGE     COMMAND                  CREATED    STATUS    PORTS    NAMES
425d6ef567d5   server    "docker-entrypoint.s..." 6 seconds ago    Up 5 seconds    0.0.0.0:8080->8080/tcp    hi_server
root@tutorial-5:~/app#

```

In the previous picture we can see how we create and run the container and how by specifying the name on the URL we get printed the greetings.



2. Analysis of the OCI / Docker image obtained from the previous step using "docker save"

I have saved the image of my server which is "server" in a tar file to analyse it by using the command "docker save -o server.tar server"

```
root@tutorial-5:~/app# docker save -o server.tar server
root@tutorial-5:~/app# ll
total 113660
drwx----- 3 scast scast      4096 May 23 10:42 ./
drwx----- 6 root  root      4096 May 23 10:28 ../
-rw-rw-r-- 1 scast scast      283 May 23 10:13 app.js
-rw-rw-r-- 1 scast scast       68 May 23 10:13 Dockerfile
-rw-rw-r-- 1 scast scast       0 May 23 10:13 .keep
drwx----- 52 scast scast     4096 May 23 10:13 node_modules/
-rw-rw-r-- 1 scast scast       53 May 23 10:13 package.json
-rw-rw-r-- 1 scast scast    31821 May 23 10:13 package-lock.json
-rw----- 1 root  root   116327424 May 23 10:42 server.tar
root@tutorial-5:~/app#
```

Now let's decompress the server.tar file and see the content with the command "tar -xf server.tar"

```
root@tutorial-5:~/app# tar -xf server.tar
root@tutorial-5:~/app# ll
total 113696
drwx----- 8 scast scast      4096 May 23 10:45 ./
drwx----- 6 root  root      4096 May 23 10:28 ../
drwxr-xr-x 2 root  root      4096 May 23 10:30 04aedabe13cf69457968e1218e6bf6874b40ff9cb15613bde5bb5f696e3fcebe/
drwxr-xr-x 2 root  root      4096 May 23 10:30 213b67ddcaaa81196a7ef59b9f45c7bfc662ee1adfab3d2607b1ce0fa631267e/
drwxr-xr-x 2 root  root      4096 May 23 10:30 66fe8211433e38a8ede87665730caf708f50c454d5fe15db272b0efa71010ecd/
drwxr-xr-x 2 root  root      4096 May 23 10:30 8792aa27a60beb94c658afc7fb54a4f1c427afe01b65cec0f2261c193ec3f495/
-rw-rw-r-- 1 scast scast      283 May 23 10:13 app.js
-rw-r--r-- 1 root  root     7118 May 23 10:30 d6be64a63ab1aaaa9206549f64408ce5a539a16437dfe2c1e34a5317ec4d2bf1.json
-rw-rw-r-- 1 scast scast       68 May 23 10:13 Dockerfile
drwxr-xr-x 2 root  root      4096 May 23 10:30 ea3e22a54d44d47cb11ba6c7337b866d5e91e231c553bb48105112b5b1637e18/
-rw-rw-r-- 1 scast scast       0 May 23 10:13 .keep
-rw-r--r-- 1 root  root     510 Jan  1 1970 manifest.json
drwx----- 52 scast scast     4096 May 23 10:13 node_modules/
-rw-rw-r-- 1 scast scast       53 May 23 10:13 package.json
-rw-rw-r-- 1 scast scast    31821 May 23 10:13 package-lock.json
-rw-r--r-- 1 root  root       89 Jan  1 1970 repositories
-rw----- 1 root  root   116327424 May 23 10:42 server.tar
```

As we can see apart from the new directories we can see two json files: "manifest.json" file and "d6be64a63ab1aaaa9206549f64408ce5a539a16437dfe2c1e34a5317ec4d2bf1.json"



```

root@tutorial-5:~/app# cat manifest.json | jq
[
  {
    "Config": "d6be64a63ab1aaaa9206549f64408ce5a539a16437dfe2c1e34a5317ec4d2bf1.json",
    "RepoTags": [
      "server:latest"
    ],
    "Layers": [
      "8792aa27a60beb94c658afc7fb54a4f1c427afe01b65cec0f2261c193ec3f495/layer.tar",
      "213b67ddcaaa81196a7ef59b9f45c7bfc662ee1adfab3d2607b1ce0fa631267e/layer.tar",
      "04aedabe13cf69457968e1218e6bf6874b40ff9cb15613bde5bb5f696e3fcebe/layer.tar",
      "ea3e22a54d44d47cb11ba6c7337b866d5e91e231c553bb48105112b5b1637e18/layer.tar",
      "66fe8211433e38a8ede87665730caf708f50c454d5fe15db272b0efa71010ecd/layer.tar"
    ]
  }
]

```

This the result of inspecting the manifest.json file.

```

root@tutorial-5:~/app# cat d6be64a63ab1aaaa9206549f64408ce5a539a16437dfe2c1e34a5317ec4d2bf1.json | jq '.history[-6:]'
[
  {
    "created": "2022-04-27T20:19:59.176076769Z",
    "created_by": "/bin/sh -c #(nop) COPY file:4d192565a7220e135cab6c77fbc1c73211b69f3d9fb37e62857b2c6eb9363d51 in /usr/local/bin/ ",
  },
  {
    "created": "2022-04-27T20:19:59.277596564Z",
    "created_by": "/bin/sh -c #(nop) ENTRYPOINT [\"docker-entrypoint.sh\"]",
    "empty_layer": true
  },
  {
    "created": "2022-04-27T20:19:59.39004664Z",
    "created_by": "/bin/sh -c #(nop) CMD [\"node\"]",
    "empty_layer": true
  },
  {
    "created": "2022-05-23T10:30:53.130890964Z",
    "created_by": "/bin/sh -c #(nop) COPY dir:a280c52462f03ee1af1bc62caef337047e299b48fd1df39bd161caf28b2dd78d in /app ",
  },
  {
    "created": "2022-05-23T10:30:53.661828694Z",
    "created_by": "/bin/sh -c #(nop) WORKDIR /app",
    "empty_layer": true
  },
  {
    "created": "2022-05-23T10:30:53.81500415Z",
    "created_by": "/bin/sh -c #(nop) CMD [\"/bin/sh\" \"-c\" \"node app.js\"]",
    "empty_layer": true
  }
]

```

Here we can see the commands which has been executed these commands correspond with the Dockerfile.



If we inspect the layer before the last layer, should match with our work directory.

```
root@tutorial-5:~/app# cd ea3e22a54d44d47cb11ba6c7337b866d5e91e231c553bb48105112b5b1637e18
root@tutorial-5:~/app/ea3e22a54d44d47cb11ba6c7337b866d5e91e231c553bb48105112b5b1637e18# ls
json layer.tar VERSION
root@tutorial-5:~/app/ea3e22a54d44d47cb11ba6c7337b866d5e91e231c553bb48105112b5b1637e18# cat json | jq
{
  "id": "ea3e22a54d44d47cb11ba6c7337b866d5e91e231c553bb48105112b5b1637e18",
  "parent": "04aedabe13cf69457968e1218e6bf6874b40ff9cb15613bde5bb5f696e3fcebe",
  "created": "1970-01-01T00:00:00Z",
  "container_config": {
    "Hostname": "",
    "Domainname": "",
    "User": "",
    "AttachStdin": false,
    "AttachStdout": false,
    "AttachStderr": false,
    "Tty": false,
    "OpenStdin": false,
    "StdinOnce": false,
    "Env": null,
    "Cmd": null,
    "Image": "",
    "Volumes": null,
    "WorkingDir": "",
    "Entrypoint": null,
    "OnBuild": null,
    "Labels": null
  },
  "os": "linux"
}
root@tutorial-5:~/app/ea3e22a54d44d47cb11ba6c7337b866d5e91e231c553bb48105112b5b1637e18# tar -xvf layer.t
ar
usr/
usr/local/
usr/local/bin/
usr/local/bin/docker-entrypoint.sh
```

```
root@tutorial-5:~/app/ea3e22a54d44d47cb11ba6c7337b866d5e91e231c553bb48105112b5b1637e18# tree .
.
├── json
├── layer.tar
├── usr
│   ├── local
│   │   └── bin
│   │       └── docker-entrypoint.sh
└── VERSION

3 directories, 4 files
```



3. Creating a system container with a distribution other than Debian and Ubuntu

I have used the alpine distribution and the variant Mini root Filesystem.

```
root@tutorial-5:/var/lib/machines# wget https://dl-cdn.alpinelinux.org/alpine/v3.15/releases/x86_64/alpine-minirootfs-3.15.4-x86_64.tar.gz
--2022-05-23 11:21:47-- https://dl-cdn.alpinelinux.org/alpine/v3.15/releases/x86_64/alpine-minirootfs-3.15.4-x86_64.tar.gz
Resolving dl-cdn.alpinelinux.org (dl-cdn.alpinelinux.org)... 199.232.18.133, 2a04:4e42:41::645
Connecting to dl-cdn.alpinelinux.org (dl-cdn.alpinelinux.org)|199.232.18.133|:443... connected.
HTTP request sent, awaiting response... 200 OK
Length: 2730061 (2.6M) [application/octet-stream]
Saving to: 'alpine-minirootfs-3.15.4-x86_64.tar.gz'

alpine-minirootfs-3.15.4- 100%[=====>] 2.60M --.-KB/s in 0.04s

2022-05-23 11:21:48 (67.7 MB/s) - 'alpine-minirootfs-3.15.4-x86_64.tar.gz' saved [2730061/2730061]
```

Here we can see the installation process and now let's see how we start the container

```
root@tutorial-5:/var/lib/machines# mkdir alpine_test01
root@tutorial-5:/var/lib/machines# tar -xf alpine-minirootfs-3.15.4-x86_64.tar.gz -C alpine_test01

root@tutorial-5:/var/lib/machines# systemd-nspawn --machine=alpinetest --directory=/var/lib/machines/alpine_test01
Spawning container alpinetest on /var/lib/machines/alpine_test01.
Press ^] three times within 1s to kill container.
alpinetest:~# cat /etc/alpine-release
3.15.4
```

```
root@tutorial-5:~# machinectl
MACHINE    CLASS    SERVICE    OS    VERSION ADDRESSES
alpinetest container systemd-nspawn alpine 3.15.4 -

1 machines listed.
```

4. Writing a unit file for Systemd, which in "systemd-analysis security" will have at least a yellow security score.

Let's create the file:

```
root@tutorial-5:/etc/systemd/system# root@tutorial-5:/etc/systemd/system# ls hi*
hi_server.service
root@tutorial-5:/etc/systemd/system#
```



```

[Unit]
Description=Hello web server wich replies "hello {name}!"according to the url.
After=network.target

[Service]
Environment=NODE_PORT=8080
Type=simple
User=root
ExecStart=/usr/bin/node /root/app/app.js
Restart=on-failure

[Install]
WantedBy=multi-user.target

```

For this service to be able to run have had to install nodejs and pm with the command “apt install nodejs” and “apt install npm”

```

root@tutorial-5:~# systemctl start hi_server.service
root@tutorial-5:~# systemctl status hi_server.service
● hi_server.service - Hello web server wich replies "hello {name}!"according to the url.
   Loaded: loaded (/etc/systemd/system/hi_server.service; disabled; vendor preset: enabled)
   Active: active (running) since Mon 2022-05-23 12:20:12 UTC; 73ms ago
     Main PID: 19724 (node)
        Tasks: 6 (limit: 4612)
       Memory: 3.8M
         CGroup: /system.slice/hi_server.service
                └─19724 /usr/bin/node /root/app/app.js

May 23 12:20:12 tutorial-5 systemd[1]: Started Hello web server wich replies "hello {name}!"according to the url.
lines 1-10/10 (END)

```

Here we have our service running, after this if we analyse the security of our service, we will see that it is unsafe.

```

root@tutorial-5:~# systemd-analyze security hi_server.service

```

NAME	DESCRIPTION
PrivateNetwork=	Service has access to the host's network
User=/DynamicUser=	Service runs as root user
CapabilityBoundingSet=~CAP_SET(UID GID PCAP)	Service may change UID/GID identities/capabilities
CapabilityBoundingSet=~CAP_SYS_ADMIN	Service has administrator privileges
CapabilityBoundingSet=~CAP_SYS_PTRACE	Service has ptrace() debugging abilities
RestrictAddressFamilies=~AF_INET INET6	Service may allocate Internet sockets
RestrictNamespaces=~CLONE_NEWUSER	Service may create user namespaces
RestrictAddressFamilies=~...	Service may allocate exotic sockets
CapabilityBoundingSet=~CAP_(CHOWN FSETID SETFCAP)	Service may change file ownership/access
CapabilityBoundingSet=~CAP_(DAC_* FOWNER IPC_OWNER)	Service may override UNIX file/IPC permissions
CapabilityBoundingSet=~CAP_NET_ADMIN	Service has network configuration privileges
CapabilityBoundingSet=~CAP_RAWIO	Service has raw I/O access
CapabilityBoundingSet=~CAP_SYS_MODULE	Service may load kernel modules
CapabilityBoundingSet=~CAP_SYS_TIME	Service processes may change the system clock
DeviceAllow=	Service has no device ACL

Let's start with the changes:

- Service run as a root user.

Let's add a user with the name of the service with the command “adduser hi_server” let's now move the app.js file which runs our server to the home directory of this new user and change the owner of the file to this user.



```

root@tutorial-5:/home/hi_server# chown hi_server:hi_server app.js
root@tutorial-5:/home/hi_server# ll
total 24
drwxr-xr-x 2 hi_server hi_server 4096 May 23 12:39 ./
drwxr-xr-x 4 root      root      4096 May 23 12:36 ../
-rw-r--r-- 1 hi_server hi_server  493 May 23 12:31 app.js

```

- Changes applied to the “hi_server.service” file.

```

[Unit]
Description=Hello web server wich replies "hello {name}!"according to the url.
After=network.target

[Service]
Environment=NODE_PORT=8080
Type=simple
User=root
ExecStart=/usr/bin/node /root/app/app.js
Restart=on-failure
CapabilityBoundingSet=
RestrictNamespaces=yes
NoNewPrivileges=true
PrivateDevices=true
PrivateMounts=true
PrivateTmp=true
PrivateUsers=true
ProtectClock=true
ProtectControlGroups=true
ProtectKernelLogs=true
ProtectKernelModules=true
ProtectKernelTunables=true
ProtectSystem=full
RestrictSUIDSGID=true
SystemCallFilter=
RestrictRealtime=true

[Install]
WantedBy=multi-user.target

```

Now let's restart the systemctl daemon with “systemct daemon-reload” and then restrart the server service.

```


root@tutorial-5:~# systemctl daemon-reload
root@tutorial-5:~# systemctl start hi_server.service
root@tutorial-5:~# systemctl status hi_server.service
● hi_server.service - Hello web server wich replies "hello {name}!"according to the url.
   Loaded: loaded (/etc/systemd/system/hi_server.service; disabled; vendor preset: enabled)
   Active: active (running) since Mon 2022-05-23 13:00:52 UTC; 4s ago
     Main PID: 20303 (node)
        Tasks: 7 (limit: 4612)
       Memory: 10.5M
      CGroup: /system.slice/hi_server.service
              └─20303 /usr/bin/node /root/app.js

May 23 13:00:52 tutorial-5 systemd[1]: Started Hello web server wich replies "hello {name}!"according to

```

Now if we inspect the security of our service we get a better score.





Hello John

```
CapabilityBoundingSet=~CAP_MAC_*           Service cannot adjust SMACK MAC
CapabilityBoundingSet=~CAP_SYS_BOOT         Service cannot issue reboot()
Delegate=                                  Service does not maintain its own delegat
LockPersonality=                           Service may change ABI personality
MemoryDenyWriteExecute=                   Service may create writable executable me
RemoveIPC=                                 Service runs as root, option does not app
RestrictNamespaces=~CLONE_NEWUTS           Service cannot create hostname namespaces
UMask=                                     Files created by service are world-readab
CapabilityBoundingSet=~CAP_LINUX_IMMUTABLE Service cannot mark files immutable
CapabilityBoundingSet=~CAP_IPC_LOCK         Service cannot lock memory into RAM
CapabilityBoundingSet=~CAP_SYS_CHROOT       Service cannot issue chroot()
ProtectHostname=                           Service may change system host/domainname
CapabilityBoundingSet=~CAP_BLOCK_SUSPEND    Service cannot establish wake locks
CapabilityBoundingSet=~CAP_LEASE            Service cannot create file leases
CapabilityBoundingSet=~CAP_SYS_PACCT        Service cannot use acct()
CapabilityBoundingSet=~CAP_SYS_TTY_CONFIG   Service cannot issue vhangup()
CapabilityBoundingSet=~CAP_WAKE_ALARM       Service cannot program timers that wake u
RestrictAddressFamilies=~AF_UNIX           Service may allocate local sockets

→ Overall exposure level for hi_server.service: 4.2 OK
```

