

# Tutorial 3


## Deployment of OpenStack

Before the exercise, create 2 VMs, one 8 GB RAM and 0.5 CPUs, the other 2 GB RAM, 0.2 CPUs. Set up SSH access without a password from the first to itself and the second.

This is the first virtual machine.

VM name <sup>?</sup>  Number of instances  ☐ Start on hold <sup>?</sup>

**Ubuntu - NI-VCC - SSD**  
69.000000 COST / HOUR



**Capacity** 19.000000 COST / HOUR

Memory <sup>?</sup>  GB

CPU <sup>?</sup>  VCPU <sup>?</sup>

**Disks** 50.000000 COST / HOUR

DISK 0: ubuntu-focal\_fossa-20220130-2040-144  GB


Custom Attributes

Heslo uzivatele root (pro pouziti ve virtualni konzoli)

This is the second virtual machine.

VM name <sup>?</sup>  Number of instances  ☐ Start on hold <sup>?</sup>

**Ubuntu - NI-VCC - SSD**  
69.000000 COST / HOUR



**Capacity** 19.000000 COST / HOUR

Memory <sup>?</sup>  GB

CPU <sup>?</sup>  VCPU <sup>?</sup>

**Disks** 50.000000 COST / HOUR

DISK 0: ubuntu-focal\_fossa-20220130-2040-144  GB

Custom Attributes

Heslo uzivatele root (pro pouziti ve virtualni konzoli)



After creating the virtual machines, we have to add an user with the name of our personal computer username in my case “scast” and then in both virtual machines we have to add this new user to the sudo group with the command “usermod -aG sudo scast”.

Now from my home and with the vpn provided by the faculty we should be able to connect by ssh with the virtual machines.

```
root@tutorial3-1: /home/scast

C:\Users\scast>ssh 10.38.6.143
The authenticity of host '10.38.6.143 (10.38.6.143)' can't be established.
ECDSA key fingerprint is SHA256:z4S2DEEMW8wr0GhFkCPRka/uKMkNugCnCes1fRDN4JM.
Are you sure you want to continue connecting (yes/no/[fingerprint])? yes
Warning: Permanently added '10.38.6.143' (ECDSA) to the list of known hosts.
scast@10.38.6.143's password:
Welcome to Ubuntu 20.04.3 LTS (GNU/Linux 5.4.0-96-generic x86_64)

 * Documentation:  https://help.ubuntu.com
 * Management:    https://landscape.canonical.com
 * Support:       https://ubuntu.com/advantage

System information as of Mon 11 Apr 2022 06:17:01 PM UTC

System load:  0.0               Processes:           111
Usage of /:   5.4% of 48.96GB   Users logged in:    1
Memory usage: 3%               IPv4 address for ens3: 10.38.6.143
Swap usage:   0%

0 updates can be applied immediately.

The list of available updates is more than a week old.
To check for new updates run: sudo apt update

The programs included with the Ubuntu system are free software;
the exact distribution terms for each program are described in the
individual files in /usr/share/doc/*/copyright.

Ubuntu comes with ABSOLUTELY NO WARRANTY, to the extent permitted by
applicable law.

To run a command as administrator (user "root"), use "sudo <command>".
See "man sudo_root" for details.

scast@tutorial3-1:~$ sudo su
[sudo] password for scast:
root@tutorial3-1:/home/scast#
```



Then I will copy the ssh keys of each machine on the other machine so they can communicate by ssh without password.

```
root@tutorial3-1:~/.ssh# root@tutorial3-1:~/.ssh# cat authorized_keys
ssh-rsa AAAAB3NzaC1yc2EAAAADAQABAAQGDAlAPN1Ph0gF6S6XtJz2obGItdMzJDpQyohQt6HyON3Xk2QyMLit/Pdp6dICNKLx10
mpjyldRIx0DnwbYt9t5k6W2Q5TwwS5BUyjtQpseI7UkjvDrOGCGGrcDMKKWYtfedJbEFw5TBCae7Yc7Z2PRCSvfHK0eOLxYwwhut58nr
Jd0N5uIb7fIhpva0sOkDcszKaxshRvg3sevxksvVJhPqTbumE1DMP+4iu9DlhwEik0HCyVAd3RgBSaF8sacpZyjdZD00ZgSFxZpf5mK
L8+G105dzGj+jhUtQUsBiM09+YBTpe2556E59As1uCEq2PX5HRU76gYY7v5JuJpnnIgNTP514n7D2IwHMTv/PPo05X9uLYmJRNufznv1
wMaVxCdkhyOFm+aeHX40u2reLqRs1wLMT627DRN1SYD69iPirJuip1R9Fsdw3nkJXQ0+Z/h9qi7vCTrbjt/HOCHG/vA10ieFzPsdZOT
10UuqK/PrE1+w7pMG4q4Ls3vKBgamP0= scast@DESKTOP-JG58M3C
ssh-rsa AAAAB3NzaC1yc2EAAAADAQABAAQGDAXtUf2Sd04bw10zD43Yjrjmfdb20p/0pZStgYMRU8evXkZf4r1ZCqd1QyIgc6bg87
nmfYJJft/BMeoA9yBW0Y2iHxJ6KSEhhEKXgZdfgWChdPyo1Uy83QXfKj73/LcK4tUXp153x1zvi3vmxFrDDQEBa77qNfY1Q2ChBjGx
oNiIGsLgetFV57DLYuBoNxyv/iJ+ZNoMXDjBoPCHc9xkR6fbvjQ/SPbHnDvWSYHcpK+FrQ7j/7n45x12RTzsvGhOnk0mJR15JdgCYM
D05FKDKYfcko1Weh2JnYnAi7GxiTMCdVvtc9GewVkf1btigm8q2U0wBiJsByZB0usp/meU2SmK6Vq6MGOM0jk5tZMCdi4KPPUwxwrvf4
VJllrI7HxH07YTCx+AAeW12kY6W85xtfZSFNYhxT1wK+jy3joZBy6Negspuf3xSVQ8kFpeQ5hK1cAQa211nS1a88Io0wDyikWNx6Ad2
X1QvNBiCyIk4en0jhwXg/R9BiQYML7E= root@tutorial3-2
ssh-rsa AAAAB3NzaC1yc2EAAAADAQABAAQGDjLsAxmOaYLkGs4e6F30ADc30dmgd+00Ye03Fm/nFxVN1d8J/ldXYPpijYGcztswx1
lxH4UMKhsuBhgafDDgDGuQnREVo2eXez69b8qCCgWJEws/i1Ic6Eum4LkoYqVDIUaebV3Zh0JdqGhKL/YE225TVs7caN0cbDVcfRYc9
ubvGKRYF9y/janCgb3bmos0YoDqm6pNuT8xk6iHUBelLQsZ0NTOfbNi4s+S9PzyIPrNxtocnbpIgzzhpE0FmDwhuVy4g+Se0+kJ8qtCH
tGb0haBPmUH26F10v/vjeZ8tt10J6Z/1UCS29XuRHsuwXo/u6qGV8+bpiXywjgyXPGs9SRjN+Bc90QeDptvGfzby1UAR6zRts2EMQjY
ITscSC56olva0yF7wfSe9GY6/jtvS5Hy6DVudqUvWozSbo/VmKXbL7dMtmfQMfYJmdRrM9FaaQm4qpk0AmBaPckjQXT3h22rLM2dhRqt
jF2UEKdWI70HKs00zG/HLJE10uSQZ8= root@tutorial3-1
```

```
root@tutorial3-2:~/.ssh# root@tutorial3-2:~/.ssh# cat authorized_keys
ssh-rsa AAAAB3NzaC1yc2EAAAADAQABAAQGDAlAPN1Ph0gF6S6XtJz2obGItdMzJDpQyohQt6HyON3Xk2QyMLit/Pdp6dICNKLx10
mpjyldRIx0DnwbYt9t5k6W2Q5TwwS5BUyjtQpseI7UkjvDrOGCGGrcDMKKWYtfedJbEFw5TBCae7Yc7Z2PRCSvfHK0eOLxYwwhut58nr
Jd0N5uIb7fIhpva0sOkDcszKaxshRvg3sevxksvVJhPqTbumE1DMP+4iu9DlhwEik0HCyVAd3RgBSaF8sacpZyjdZD00ZgSFxZpf5mK
L8+G105dzGj+jhUtQUsBiM09+YBTpe2556E59As1uCEq2PX5HRU76gYY7v5JuJpnnIgNTP514n7D2IwHMTv/PPo05X9uLYmJRNufznv1
wMaVxCdkhyOFm+aeHX40u2reLqRs1wLMT627DRN1SYD69iPirJuip1R9Fsdw3nkJXQ0+Z/h9qi7vCTrbjt/HOCHG/vA10ieFzPsdZOT
10UuqK/PrE1+w7pMG4q4Ls3vKBgamP0= scast@DESKTOP-JG58M3C
ssh-rsa AAAAB3NzaC1yc2EAAAADAQABAAQGDjLsAxmOaYLkGs4e6F30ADc30dmgd+00Ye03Fm/nFxVN1d8J/ldXYPpijYGcztswx1
lxH4UMKhsuBhgafDDgDGuQnREVo2eXez69b8qCCgWJEws/i1Ic6Eum4LkoYqVDIUaebV3Zh0JdqGhKL/YE225TVs7caN0cbDVcfRYc9
ubvGKRYF9y/janCgb3bmos0YoDqm6pNuT8xk6iHUBelLQsZ0NTOfbNi4s+S9PzyIPrNxtocnbpIgzzhpE0FmDwhuVy4g+Se0+kJ8qtCH
tGb0haBPmUH26F10v/vjeZ8tt10J6Z/1UCS29XuRHsuwXo/u6qGV8+bpiXywjgyXPGs9SRjN+Bc90QeDptvGfzby1UAR6zRts2EMQjY
ITscSC56olva0yF7wfSe9GY6/jtvS5Hy6DVudqUvWozSbo/VmKXbL7dMtmfQMfYJmdRrM9FaaQm4qpk0AmBaPckjQXT3h22rLM2dhRqt
jF2UEKdWI70HKs00zG/HLJE10uSQZ8= root@tutorial3-1
```

Now we will continue by deleting firewall because it breaks open stack installation by executing the command “apt purge firewall” in both virtual machines.

```
root@tutorial3-1:~# apt purge firewall
Reading package lists... Done
Building dependency tree
Reading state information... Done
The following packages were automatically installed and are no longer required:
  gir1.2-nm-1.0 ipset libipset13 libjansson4 libnftables1 libnm0 python3-decorator python3-firewall
  python3-nftables python3-selinux python3-slip python3-slip-dbus
Use 'sudo apt autoremove' to remove them.
```

```
root@tutorial3-2:~# apt purge firewall
Reading package lists... Done
Building dependency tree
Reading state information... Done
The following packages were automatically installed and are no longer required:
  gir1.2-nm-1.0 ipset libipset13 libjansson4 libnftables1 libnm0 python3-decorator python3-firewall
```

(BOTH MACHINES)

Now we are going to start installing open stack with Kolla. First, we will install dependencies, for that we will run the next commands “apt update”, “apt install python3-dev libffi-dev gcc libssl-dev” and “apt install python3-pip” in both virtual machines, after installing pip we need to make sure that we upgrade pip to the latest version executing “pip3 install -U pip”.



```

Processing triggers for libc-bin (2.31-0ubuntu9.2) ...
root@tutorial3-1:~# pip3 install -U pip
Collecting pip
  Downloading pip-22.0.4-py3-none-any.whl (2.1 MB)
    | 2.1 MB 9.4 MB/s
Installing collected packages: pip
  Attempting uninstall: pip
    Found existing installation: pip 20.0.2
    Not uninstalling pip at /usr/lib/python3/dist-packages, outside environment /usr
    Can't uninstall 'pip'. No files were found to uninstall.
Successfully installed pip-22.0.4

```

The next step is to install ansible but we need to install it from pip by executing “`pip install -U 'ansible<3.0'`” only on the first virtual machine.

```

Successfully installed pip-22.0.4
root@tutorial3-1:~# pip install -U 'ansible<3.0'
Collecting ansible<3.0
  Downloading ansible-2.10.7.tar.gz (29.9 MB)
    29.9/29.9 MB 20.5 MB/s eta 0:00:00
  Preparing metadata (setup.py) ... done
Collecting ansible-base<2.11,>=2.10.5
  Downloading ansible-base-2.10.17.tar.gz (6.1 MB)
    6.1/6.1 MB 38.2 MB/s eta 0:00:00
  Preparing metadata (setup.py) ... done
Requirement already satisfied: PyYAML in /usr/lib/python3/dist-packages (from ansible-base<2.11,>=2.10.5->ansible<3.0) (5.3.1)

```

Now, we have to install Kolla ansible by executing “`pip3 install git+https://opendev.org/openstack/kolla-ansible@stable/xena`”, then we have to create kola directory with “`mkdir -p /etc/kolla`” then, copy `globals.yml` and `passwords.yml` with “`cp -r /usr/local/share/kolla-ansible/etc_examples/kolla/* /etc/kolla`”, then copy `all-in-one` and `multinode` inventory files to the current directory with “`cp /usr/local/share/kolla-ansible/ansible/inventory/* .`”

The next part is going to be to edit the configuration files,”`mkdir /etc/ansible/`” for this we will edit “`/etc/ansible/ansible.cfg`”

```

root@tutorial3-1: ~
[defaults]
host_key_checking=False
pipelining=True
forks=100

```

Now we need to customize the multimode file with the next aspects.



```
# These initial groups are the only groups required to be modified. The
# additional groups are for more control of the environment.
[control]
# These hostname must be resolvable from your deployment host
10.38.6.143

# The above can also be specified as follows:
#control[01:03]    ansible_user=kolla

# The network nodes are where your l3-agent and loadbalancers will run
# This can be the same as a host in the control group
[network:children]
control

[compute]
10.38.6.149

[monitoring:children]
control

# When compute nodes and control nodes use different interfaces,
# you need to comment out "api_interface" and other interfaces from the globals.yml
# and specify like below:
#compute01 neutron_external_interface=eth0 api_interface=em1 storage_interface=em1 tunnel_interface=em1

[storage:children]
compute
```

If we try now to ping with the command “`ansible -i multinode all -m ping`” we will see these three successes.

```
root@tutorial3-1:~# ansible -i multinode all -m ping
[WARNING]: Invalid characters were found in group names but not replaced, use -vvvv to see details
[DEPRECATION WARNING]: Distribution Ubuntu 20.04 on host localhost should use /usr/bin/python3, but is
using /usr/bin/python for backward compatibility with prior Ansible releases. A future Ansible release
will default to using the discovered platform python for this host. See
https://docs.ansible.com/ansible/2.10/reference_appendices/interpreter_discovery.html for more
information. This feature will be removed in version 2.12. Deprecation warnings can be disabled by
setting deprecation_warnings=False in ansible.cfg.
localhost | SUCCESS => {
  "ansible_facts": {
    "discovered_interpreter_python": "/usr/bin/python"
  },
  "changed": false,
  "ping": "pong"
}
[DEPRECATION WARNING]: Distribution Ubuntu 20.04 on host 10.38.6.149 should use /usr/bin/python3, but
is using /usr/bin/python for backward compatibility with prior Ansible releases. A future Ansible
release will default to using the discovered platform python for this host. See
https://docs.ansible.com/ansible/2.10/reference_appendices/interpreter_discovery.html for more
information. This feature will be removed in version 2.12. Deprecation warnings can be disabled by
setting deprecation_warnings=False in ansible.cfg.
10.38.6.149 | SUCCESS => {
  "ansible_facts": {
    "discovered_interpreter_python": "/usr/bin/python"
  },
  "changed": false,
  "ping": "pong"
}
[DEPRECATION WARNING]: Distribution Ubuntu 20.04 on host 10.38.6.143 should use /usr/bin/python3, but
is using /usr/bin/python for backward compatibility with prior Ansible releases. A future Ansible
release will default to using the discovered platform python for this host. See
https://docs.ansible.com/ansible/2.10/reference_appendices/interpreter_discovery.html for more
information. This feature will be removed in version 2.12. Deprecation warnings can be disabled by
setting deprecation_warnings=False in ansible.cfg.
10.38.6.143 | SUCCESS => {
  "ansible_facts": {
    "discovered_interpreter_python": "/usr/bin/python"
  },
  "changed": false,
  "ping": "pong"
}
```

The next step will be to run Kolla first we have to execute “`kolla-genpwd`”, after this we will open the file `/etc/kolla/globals.yml` and edit the next lines:

```
# Valid options are ['centos', 'debian', 'rhel', 'ubuntu']
kolla_base_distro: "ubuntu"
```



```
# 'network_interface' as set in the Networking section below.  
kolla_internal_vip_address: "10.38.6.143"
```

This is the ip address of the first virtual machine

```
# followed for other types of interfaces.  
network_interface: "ens3"
```

```
# addresses for that reason.  
neutron_external_interface: "veth0"
```

```
#enable_hacluster: "no"  
enable_haproxy: "no"
```

Now we have to create virtual interface 0 for internet, so let's open the file /etc/network/interfaces on both virtual machines and apply the same changes

```
auto lo  
iface lo inet loopback  
  
auto ens3  
iface ens3 inet static  
    pre-up ip link add veth0 type veth peer name veth1  
    pre-up ip link set veth0 up  
    pre-up ip link set veth1 up  
    address 10.38.6.208  
    network 10.38.4.0  
    netmask 255.255.252.0  
    gateway 10.38.4.1  
  
source /etc/network/interfaces.d/*.cfg  
~
```

After this we must restart networking service by executing “service networking restart”.

Then to check if everything is correct, we will execute the command “ip link” and we should get a response like this.

```
root@tutorial3-2:~# ip link  
1: lo: <LOOPBACK,UP,LOWER_UP> mtu 65536 qdisc noqueue state UNKNOWN mode DEFAULT group default qlen 1000  
    link/loopback 00:00:00:00:00:00 brd 00:00:00:00:00:00  
2: ens3: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc fq_codel state UP mode DEFAULT group default qlen 1000  
    link/ether 02:00:0a:26:06:95 brd ff:ff:ff:ff:ff:ff  
3: veth1@veth0: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc noqueue state UP mode DEFAULT group default qlen 1000  
    link/ether 0a:a4:54:2c:8c:b4 brd ff:ff:ff:ff:ff:ff  
4: veth0@veth1: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc noqueue state UP mode DEFAULT group default qlen 1000  
    link/ether be:91:c4:b1:b3:96 brd ff:ff:ff:ff:ff:ff  
root@tutorial3-2:~#
```

After configuration is set, we can proceed to the deployment phase. Now back in the first vm (ctl) we need to setup basic host-level dependencies we have to run “kolla-ansible -i ./multinode bootstrap-servers -e ansible\_python\_interpreter=/usr/bin/python3” with the flag -e to specify the python interpreter.





```

skipping: [10.38.6.143]
skipping: [10.38.6.149]

TASK [baremetal : Update apt cache] *****
*****
***
changed: [10.38.6.143]
changed: [10.38.6.149]

TASK [baremetal : Set firewall default policy] *****
*****
***
ok: [10.38.6.143]
ok: [10.38.6.149]

```

After bootstrap the next is prechecks so let's execute "kolla-ansible -i ./multinode prechecks -e ansible\_python\_interpreter=/usr/bin/python3"

```

PLAY [Apply role masakari] *****
skipping: no hosts matched

PLAY RECAP *****
10.38.6.143      : ok=68  changed=0    unreachable=0    failed=0    skipped=60    rescued=0
ignored=0
10.38.6.149      : ok=30  changed=0    unreachable=0    failed=0    skipped=30    rescued=0
ignored=0
localhost       : ok=11  changed=0    unreachable=0    failed=0    skipped=14    rescued=0
ignored=0

```

And the last command "kolla-ansible -i ./multinode deploy -e ansible\_python\_interpreter=/usr/bin/python3"

```

TASK [common : Creating log volume] *****
changed: [10.38.6.149]
changed: [10.38.6.143]

TASK [common : Link kolla_logs volume to /var/log/kolla] *****
changed: [10.38.6.149]
changed: [10.38.6.143]

RUNNING HANDLER [common : Restart fluentd container] *****
changed: [10.38.6.149]
changed: [10.38.6.143]

RUNNING HANDLER [common : Restart kolla-toolbox container] *****
changed: [10.38.6.149]
changed: [10.38.6.143]

```

If we ssh on a second terminal and execute "docker ps" we can see that we are creating the containers for the database

```

root@tutorial13-1:~# docker ps

```

CONTAINER ID	IMAGE	PORTS	NAMES	COMMAND	CREATED
d8836a87929a	quay.io/openstack.kolla/ubuntu-source-mariadb-server:xena		mariadb	"dumb-init -- kolla_..."	7 seconds ago
8be1f231592c	quay.io/openstack.kolla/ubuntu-source-cron:xena		cron	"dumb-init --single-..."	45 seconds ago
c037105299ff	quay.io/openstack.kolla/ubuntu-source-kolla-toolbox:xena		kolla_toolbox	"dumb-init --single-..."	About a minute ago
27bdc2adba31	quay.io/openstack.kolla/ubuntu-source-fluentd:xena		fluentd	"dumb-init --single-..."	About a minute ago

```

root@tutorial13-1:~#

```

Now when the deploy is finished, let's install the openstack client with the command "apt install python3-openstackclient" and after that "kolla-ansible post-deploy" this will generate the file /etc/kolla/admin\_openrc.sh



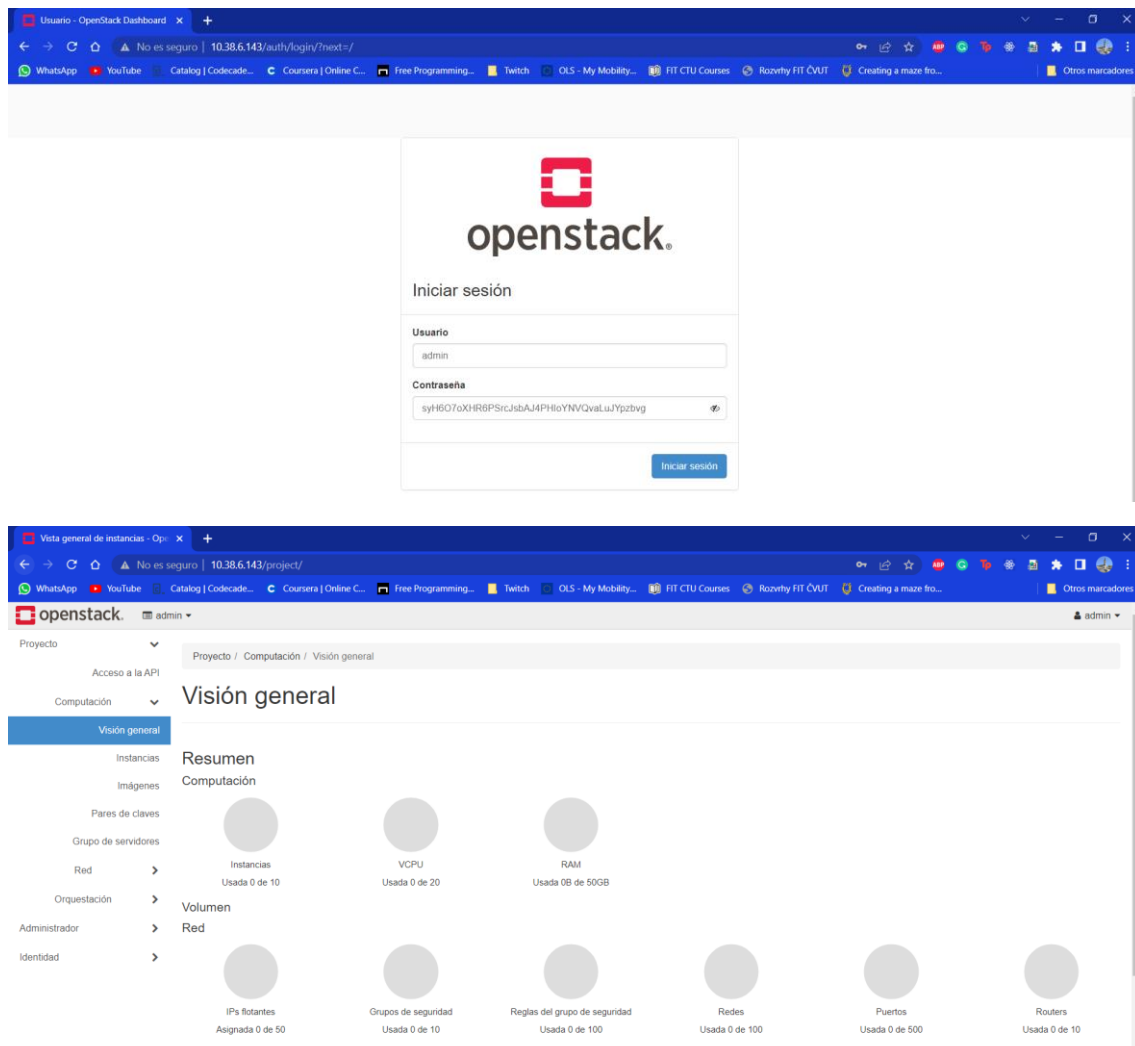
```

root@tutorial13-1:~# cat /etc/kolla/admin-openrc.sh
# Ansible managed

# Clear any old environment that may conflict.
for key in $( set | awk '{FS="="} / ^OS_ / {print $1}' ); do unset $key ; done
export OS_PROJECT_DOMAIN_NAME=Default
export OS_USER_DOMAIN_NAME=Default
export OS_PROJECT_NAME=admin
export OS_TENANT_NAME=admin
export OS_USERNAME=admin
export OS_PASSWORD=syH607oXHR6PSrcJsbAJ4PHIoYNNQvaLuJYpzbvg
export OS_AUTH_URL=http://10.38.6.143:35357/v3
export OS_INTERFACE=internal
export OS_ENDPOINT_TYPE=internalURL
export OS_IDENTITY_API_VERSION=3
export OS_REGION_NAME=RegionOne
export OS_AUTH_PLUGIN=password
root@tutorial13-1:~#

```

So now if we try to connect to the highlighted ip address via our browser we should be able to log in openstack graphical interface with user “admin” and the password is the one provided by the previous picture.



And then the last command “/usr/local/share/kolla-ansible/init-runonce”





```
Done.

To deploy a demo instance, run:

openstack server create \
  --image cirros \
  --flavor m1.tiny \
  --key-name mykey \
  --network demo-net \
  demo1
root@tutorial3-1:~# openstack server create \
> --image cirros \
> --flavor m1.tiny \
> --key-name mykey \
> --network demo-net \
> demo1
```

When the script will be done it will show us how to deploy a demo instance of the machine and we will execute as is shown in the previous image.

Finally in the graphical interface we will have access to this instance.

The screenshot displays the OpenStack dashboard interface. The top navigation bar shows the 'openstack.' logo and the user 'admin'. The left sidebar contains a navigation menu with options like Project, API Access, Compute, Overview, Instances, Images, Key Pairs, Server Groups, Network, and Orchestration. The main content area is titled 'Instances' and shows a table with one instance named 'demo1'. The instance is in the 'Active' state, using the 'cirros' image, 'm1.tiny' flavor, and 'mykey' key pair. Below the table, there is a console view for the instance 'demo1'. The console output shows the boot process of the instance, including messages about freeing kernel image memory, checking W+X mappings, and detecting a new USB device (QEMU USB Tablet). The console also shows the login prompt for the 'cirros' user with the default password 'gocubsgo'.

Instance Name	Image Name	IP Address	Flavor	Key Pair	Status	Availability Zone	Task	Power State	Age	Actions
demo1	cirros	10.0.0.124	m1.tiny	mykey	Active	nova	None	Running	2 minutes	Create Snapshot



Exercise 3: OpenStack which can run VMs on two nodes (either a second computing or compute service on the control node). The VMs must be able to ping each other.

Let's edit the multimode file and on the [compute] group let's add the ip address of the control node too.

```
root@tutorial3-ctl: ~
# These initial groups are the only groups required to be modified. The
# additional groups are for more control of the environment.
[control]
# These hostname must be resolvable from your deployment host
10.38.6.69

# The above can also be specified as follows:
#control[01:03]    ansible_user=kolla

# The network nodes are where your l3-agent and loadbalancers will run
# This can be the same as a host in the control group
[network:children]
control

[compute]
10.38.6.69
10.38.6.81
```

Now we have to make update the prechecks so the changes are reflected and the launch the deploy command for openstack again, that's why we will execute "kolla-ansible -i ./multinode prechecks -e ansible\_python\_interpreter=/usr/bin/python3"

```
PLAY [Apply role masakari] *****skipping: no hosts matched

PLAY RECAP *****10.38.6.69 : 0
k=50 changed=0 unreachable=0 failed=0 skipped=82 rescued=0 ignored=0
10.38.6.81 : ok=25 changed=0 unreachable=0 failed=0 skipped=31 rescued=0 ignored=0
localhost : ok=11 changed=0 unreachable=0 failed=0 skipped=14 rescued=0 ignored=0
```

and then "kolla-ansible -i ./multinode deploy -e ansible\_python\_interpreter=/usr/bin/python3"

```
PLAY [Apply role masakari] *****skipping: no hosts matched

PLAY RECAP *****10.38.6.69 : ok=235 changed=29 unreachable=0 failed=0 skipped=123 rescued=0 ignored=0
10.38.6.81 : ok=49 changed=1 unreachable=0 failed=0 skipped=43 rescued=0 ignored=0
localhost : ok=4 changed=0 unreachable=0 failed=0 skipped=0 rescued=0 ignored=0

root@tutorial3-ctl:~# kolla-ansible post-deploy
Post-Deploying Playbooks : ansible-playbook -e @/etc/kolla/globals.yml -e @/etc/kolla/passwords.yml -e CONFIG_DIR=/etc/kolla /usr/local/share/kolla-ansible/ansible/post-deploy.yml
[WARNING]: no inventory was parsed, only implicit localhost is available
[WARNING]: provided hosts list is empty, only localhost is available. Note that the implicit localhost does not match 'all'

PLAY [Creating admin openrc file on the deploy node] *****
TASK [Gathering Facts] *****ok: [localhost]
TASK [Template out admin-openrc.sh] *****ok: [localhost]
TASK [octavia : Template out octavia-openrc.sh] *****skipping: [localhost]

PLAY RECAP *****localhost : ok=2 changed=0 unreachable=0 failed=0 skipped=1 rescued=0 ignored=0
```

Now if we execute the command "openstack host list" we will see both virtual machines as compute nodes.



```

root@tutorial3-ctl:~# openstack host list
+-----+-----+-----+
| Host Name      | Service  | Zone    |
+-----+-----+-----+
| tutorial3-ctl  | conductor| internal|
| tutorial3-cmp  | compute  | nova    |
| tutorial3-ctl  | compute  | nova    |
| tutorial3-ctl  | scheduler| internal|
+-----+-----+-----+

```

Now we will create a server specifying the compute node we want to use with the command “openstack server create \

```

--image cirros \
--flavor m1.tiny \
--key-name mykey \
--network demo-net \
--availability-zone nova:tutorial3-ctl \
demo2”

```

```

root@tutorial3-ctl:~# openstack server create \
> --image cirros \
> --flavor m1.tiny \
> --key-name mykey \
> --network demo-net \
> --availability-zone nova:tutorial3-ctl \
> demo2
+-----+-----+
| Field                                | Value |
+-----+-----+
| OS-DCF:diskConfig                    | MANUAL |
| OS-EXT-AZ:availability_zone          | nova   |
| OS-EXT-SRV-ATTR:host                 | None   |
| OS-EXT-SRV-ATTR:hypervisor_hostname | None   |
| OS-EXT-SRV-ATTR:instance_name        |        |
| OS-EXT-STS:power_state                | NOSTATE |
| OS-EXT-STS:task_state                 | scheduling |
| OS-EXT-STS:vm_state                   | building |
| OS-SRV-USG:launched_at                | None   |
| OS-SRV-USG:terminated_at              | None   |
+-----+-----+

```

Finally, to check that everything has been made correctly if in the GUI we check the nodes of the instances we will see that they have different nodes, and they can ping between each other across the network.



openstack.

admin

Project

API Access

Compute

Overview

Instances

Images

Key Pairs

Server Groups

Network

Orchestration

Admin

Project / Compute / Instances

Instances

Instance ID =

Mostrando 2 artículos

	Instance Name	Image Name	IP Address	Flavor	Key Pair	Status		Availability Zone
<input type="checkbox"/>	demo2	cirros	10.0.0.3	m1.tiny	mykey	Active		nova
<input type="checkbox"/>	demo1	cirros	10.0.0.219	m1.tiny	mykey	Active		nova

Mostrando 2 artículos

openstack.

admin

Project

API Access

Compute

Overview

Instances

Images

Key Pairs

Server Groups

Network

Orchestration

Admin

Identity

Project / Compute / Instances / demo1

demo1

Overview

Interfaces

Log

Console

Action Log

Name

ID

Description

Project ID

Status

Locked

Availability Zone

Created

Age

Host

Instance Name

Reservation ID

Launch Index

Hostname

Kernel ID

Ramdisk ID

Device Name

User Data

demo1

26f4a398-0eff-42f5-bb7d-4b37f3f981

demo1

d5ce5fd20895453ea086356555ca5b9e

Active

False

nova

May 8, 2022, 4:41 p.m.

18 hours, 24 minutes

tutorial3-cmp

instance-00000001

r-ble040us

-

demo1

-

-

/dev/vda

-

Specs



openstack.

admin

Project

API Access

Compute

Overview

Instances

Images

Key Pairs

Server Groups

Network

Orchestration

Admin

Identity

Project / Compute / Instances / demo2

demo2

Overview

Interfaces

Log

Console

Action Log

Name

 demo2

ID

 4553c346-1098-48c5-9b27-55e2291b0a17

Description

 demo2

Project ID

 d5ce5fd20895453ea086356555ca5b9e

Status

 Active

Locked

 False

Availability Zone

 nova

Created

 May 9, 2022, 11:02 a.m.

Age

 3 minutes

Host

 tutorial3-ctl

Instance Name

 instance-00000006

Reservation ID

 r-72cbosjb

Launch Index

 -

Hostname

 demo2

Kernel ID

 -

Ramdisk ID

 -

Device Name

 /dev/vda

User Data

 -

Specs

openstack.

admin

Orchestration

Admin

Identity

```

further output written to /dev/ttyS0
[ 3.021445] virtio_blk virtio1: [vdal 2097152 512-byte logical blocks (1.07 GiB/1.09 GiB)]
[ 3.043662] GPT:Primary header thinks Alt. header is not at the end of the disk.
[ 3.059630] GPT:229375 != 2097151
[ 3.066167] GPT:Alternate GPT header not at the end of the disk.
[ 3.070662] GPT:229375 != 2097151
[ 3.086504] GPT: Use GNU Parted to correct GPT errors.

login as 'cirros' user, default password: 'g0cubsg0', use 'sudo' for root.
demo2 login: cirros
Password:
$ ping 10.0.0.219
PING 10.0.0.219 (10.0.0.219): 56 data bytes
64 bytes from 10.0.0.219: seq=0 ttl=64 time=5.708 ms
64 bytes from 10.0.0.219: seq=1 ttl=64 time=5.146 ms
64 bytes from 10.0.0.219: seq=2 ttl=64 time=2.333 ms
64 bytes from 10.0.0.219: seq=3 ttl=64 time=2.107 ms
^C
--- 10.0.0.219 ping statistics ---
4 packets transmitted, 4 packets received, 0% packet loss
round-trip min/avg/max = 2.107/3.843/5.708 ms
$

```



Exercise 4: As in the previous point, plus functional access from the VM to the Internet using NAT.

First, I will allow nat on iptables of the network node with the next rules:

```
iptables -t nat -A POSTROUTING -o ens3 -j MASQUERADE
```

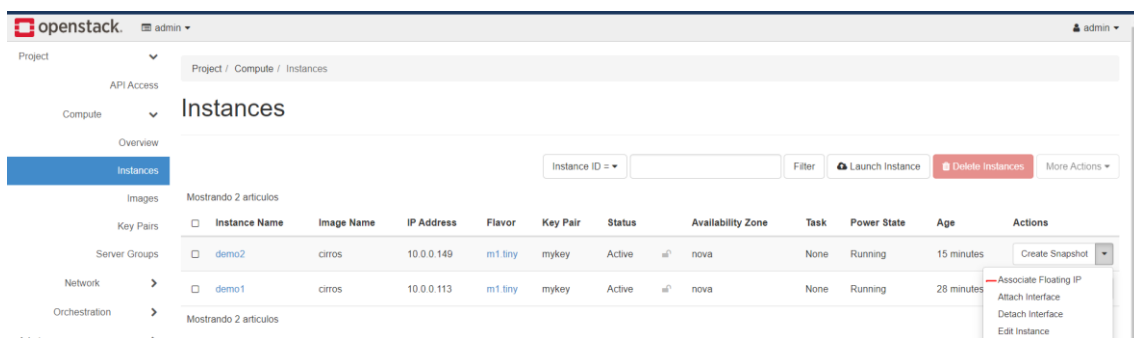
```
root@tutorial3-ctl:~# iptables -t nat -A POSTROUTING -o ens3 -j MASQUERADE
root@tutorial3-ctl:~# iptables -L -t nat
Chain PREROUTING (policy ACCEPT)
target     prot opt source               destination

Chain INPUT (policy ACCEPT)
target     prot opt source               destination

Chain OUTPUT (policy ACCEPT)
target     prot opt source               destination

Chain POSTROUTING (policy ACCEPT)
target     prot opt source               destination
MASQUERADE all  --  anywhere             anywhere
```

Then we need to add floating ip to each instance.



Instance Name	Image Name	IP Address	Flavor	Key Pair	Status	Availability Zone	Task	Power State	Age	Actions
demo2	cirros	10.0.0.149	m1.tiny	mykey	Active	us-east-1a	None	Running	15 minutes	Create Snapshot
demo1	cirros	10.0.0.113	m1.tiny	mykey	Active	us-east-1a	None	Running	28 minutes	Associate Floating IP, Attach Interface, Detach Interface, Edit Instance

### Manage Floating IP Associations

**IP Address \***

No floating IP addresses allocated +

Select the IP address you wish to associate with the selected instance or port.

**Port to be associated \***

demo2: 10.0.0.149

Cancel Associate



## Allocate Floating IP

Pool \*

public1

Description

Description:

Allocate a floating IP from a given floating IP pool.

Project Quotas

Floating IP

0 of 50 Used

Cancel

Allocate IP

## Manage Floating IP Associations

IP Address \*

10.0.2.186

+

Port to be associated \*

demo2: 10.0.0.149

Select the IP address you wish to associate with the selected instance or port.

Cancel

Associate

Add the ip of the gateway of the external subnet to the interface veth1 with

“ip addr add 10.0.2.1/24 dev veth1”

```
root@tuorial3-ctl:~# ip addr add 10.0.2.1/24 dev veth1
```

And finally, to activate the forwarding of packages in the control node let's execute “echo 1 > /proc/sys/net/ipv4/ip\_forward”

```
root@tuorial3-ctl:~# echo 1 > /proc/sys/net/ipv4/ip_forward
```

And now if we ping from one of our instances to google for example (8.8.8.8) we should get a reply.





```
64 bytes from 8.8.8.8: seq=1 ttl=56 time=2.083 ms
64 bytes from 8.8.8.8: seq=2 ttl=56 time=2.359 ms
64 bytes from 8.8.8.8: seq=3 ttl=56 time=3.600 ms
64 bytes from 8.8.8.8: seq=4 ttl=56 time=2.245 ms
64 bytes from 8.8.8.8: seq=5 ttl=56 time=2.194 ms
64 bytes from 8.8.8.8: seq=6 ttl=56 time=2.439 ms
64 bytes from 8.8.8.8: seq=7 ttl=56 time=2.545 ms
64 bytes from 8.8.8.8: seq=8 ttl=56 time=2.431 ms
64 bytes from 8.8.8.8: seq=9 ttl=56 time=1.941 ms
64 bytes from 8.8.8.8: seq=10 ttl=56 time=3.058 ms
^C
--- 8.8.8.8 ping statistics ---
11 packets transmitted, 11 packets received, 0% packet loss
round-trip min/avg/max = 1.941/2.594/3.646 ms
$ ping www.google.com
PING www.google.com (142.251.36.100): 56 data bytes
64 bytes from 142.251.36.100: seq=0 ttl=56 time=1.837 ms
64 bytes from 142.251.36.100: seq=1 ttl=56 time=3.194 ms
64 bytes from 142.251.36.100: seq=2 ttl=56 time=2.717 ms
64 bytes from 142.251.36.100: seq=3 ttl=56 time=3.106 ms
^C
--- www.google.com ping statistics ---
4 packets transmitted, 4 packets received, 0% packet loss
round-trip min/avg/max = 1.837/2.713/3.194 ms
$ _
```

